

# Project 2: User Programs

## Preliminaries

Fill in your name and email address.

Sun Shaocong [2000012977@stu.pku.edu.cn](mailto:2000012977@stu.pku.edu.cn)

If you have any preliminary comments on your submission, notes for the TAs, please give them here.

None.

Please cite any offline or online sources you consulted while preparing your submission, other than the Pintos documentation, course text, lecture notes, and course staff.

About overall architecture of system calls:

- <https://github.com/mit-pdos/xv6-riscv/blob/riscv/kernel/syscall.c>

About files struct of a thread:

- <https://blog.csdn.net/mediatec/article/details/86735691>

## Argument Passing

### DATA STRUCTURES

A1: Copy here the declaration of each new or changed struct or struct member, global or static variable, typedef, or enumeration. Identify the purpose of each in 25 words or less.

Definations:

```
#define ALIGNMENT 4                /* Alignment. */
#define DECREASE_PTR(ptr,bytes)    /* Decrease PTR by BYTES. */
#define PUT_4BYTES(ptr,value)     /* Put a 4 bytes VALUE at PTR. */
#define ALIGN(ptr)                /* Round PTR down to a multiple of ALIGNMENT. */
#define OVERFLOW(cur_esp,init_esp) /* Stack overflow? */
#define TEST_AND_PUSH(ptr,value,init_esp)
/* Push a 4 bytes VALUE on stack PTR, but check before push. Return whether push
succeeds. */
```

### ALGORITHMS

A2: Briefly describe how you implemented argument parsing. How do you arrange for the elements of `argv[]` to be in the right order? How do you avoid overflowing the stack page?

- If a command line has B bytes, the total length of all arguments(include '\0') is at most B+1. So, after a return from load, I save initial `if_esp` as `save_esp` and decrease it by B+1 to leave enough space for the

argument strings. Use `save_esp` to push argument strings while pushing the address of each string through `if_esp`. Make alignment and store a NULL pointer at `if_esp`. Each time we get a token out of command line, decrease `save_esp` by its length(include `'\0'`) and copy the token string to `save_esp`, after that push the value of `save_esp`(string address) at `if_esp`. When all tokens are processed, argument strings are copied to user stack and their addresses are stored in a left-to-right order. `Argv` is current `if_esp` and `argc` can be counted before. To make the elements of `argv[]` in a right order, just **reverse the argv array**. The rest is relatively simple. Push `argv`, `argc`, "return address" at `if_esp` and finally done argument passing.

- Another thing is about how to avoid overflowing the stack page. As you may notice, I define `TEST_AND_PUSH` rather than `PUSH`. This inline routine will test stack overflow before truly putting the value to stack. I ensure that whenever to put a value at address A, a lower address B shall be tested some time earlier. So all writing is safe.

## RATIONALE

A3: Why does Pintos implement `strtok_r()` but not `strtok()`?

- Maybe because `strtok_r()` is more flexible in use. The former one returns `save_ptr` to the user. The latter one only saves `save_ptr` in a fixed location. After one call to `strtok()` on string S, we have to finish all calls on string S because we don't hold any pointer to the rest part of string S. If a call to `strtok()` on string T break in, we will lose the opportunity to operate on the rest of S.

A4: In Pintos, the kernel separates commands into a executable name and arguments. In Unix-like systems, the shell does this separation. Identify at least two advantages of the Unix approach.

- It will decrease kernel processing time.
- The shell can make some validity check before passing command to kernel, reduce the possibility of kernel errors.

## System Calls

### DATA STRUCTURES

B1: Copy here the declaration of each new or changed struct or struct member, global or static variable, typedef, or enumeration. Identify the purpose of each in 25 words or less.

New struct member in struct `thread`:

```
/* Used in filesys syscall. */
struct files_struct files;          /**< Files struct. */
/* Used in syscall exec, wait and exit. */
struct thread *parent;             /**< Current thread's parent. */
/* Used in syscall exec to synchronize during loading. */
bool child_load_success;           /**< Child loading result. */
struct semaphore child_load_sema;  /**< Semaphore for load. */
/* Used in syscall exit and wait. */
bool is_parent_died;               /**< Its parent died or not. */
```

```

struct list child_list;           /**< A list of its child. */
struct child *child_msg;          /**< Child struct of this thread
                                  owned by its parent. */

/* Used for deny write to executables. */
struct file *executable;          /**< Executable file. */

```

New definitions in thread.h:

```

/** Thread's open file definition. */
typedef int fd_t;                  /**< File descriptor. */
#define FD_ERROR ((fd_t) -1)      /**< Error value for fd_t. */
typedef int fd_set_t;             /**< Bitmap of opening fds. */
#define INIT_FD_SET 0x3           /**< Initial fd set. */
#define MAX_FILE 8*sizeof(fd_set_t) /**< Maximum files a thread can open. */
#define USER_FD_MIN (fd_t)2       /**< Lowest fd user can open. */
#define USER_FD_MAX (fd_t)(MAX_FILE-1) /**< Highest fd user can open. */
#define STDIN_FD 0                 /**< STDIN file descriptor. */
#define STDOUT_FD 1                /**< STDOUT file descriptor. */
/** Three basic operation on a fd set. */
#define IS_FREE_FD(fd_set,fd)      /* Is FD allocatable in FD_SET. */
#define ALLOCATE_FD(fd_set,fd)     /* Add FD to FD_SET. */
#define REMOVE_FD(fd_set,fd)       /* Remove FD from FD_SET. */

```

New structs in thread.h:

```

/** A thread's opening files struct. */
struct files_struct
{
    fd_set_t      fd_set;          /**< Open fd set. */
    struct file    *fd_table[MAX_FILE]; /**< Open fd table. */
};

/** A child struct of a thread. */
struct child
{
    struct thread *t;              /**< Child thread struct. */
    tid_t tid;                     /**< Child tid. */
    bool is_terminated;            /**< Child exit or not. */
    int saved_exit_status;         /**< Saved exit status. */
    struct semaphore terminated_sema; /**< Semaphore for wait. */
    struct list_elem elem;         /**< List elem. */
};

```

New definitions and variables in syscall.c:

```

#define SYSNUM 13
#define MAXARGS 4

static bool in_syscall_context;          /**< Are we processing a syscall? */

static struct lock filesys_lock;         /**< A lock for filesys operations. */
static struct lock exit_lock;           /**< A lock for exit call. */

static int (*syscalls[])(uint32_t *argv); /** Handler list. */
static int _argc[];                     /** Argc list. */

```

B2: Describe how file descriptors are associated with open files. Are file descriptors unique within the entire OS or just within a single process?

- An opening file descriptor is recorded in `fd_set` as a bit "1" for a thread. Specifically, if file descriptor FD is opened, The FD bit of `fd_set` in thread's files struct will be set to 1 and `fd_table[FD]` will be pointed to file struct of the open file.
- File descriptors are unique just within a single process.

Following functions are provided for manipulate files struct of a single thread in `thread.h`:

```

fd_t thread_open_file(struct file*);
int thread_close_file(fd_t fd);
void thread_close_all(void);
struct file *thread_get_file(fd_t fd);

```

## ALGORITHMS

B3: Describe your code for reading and writing user data from the kernel.

- Reading and writing are implemented based on `get_user()` and `put_user()` provided in the Pintos documentation. To make sure these two functions work correctly, `page_fault()` is also modified according to the Pintos documentation. When page fault detected during system call processing, page fault handler returns control back to syscall handler instead of killing the process. Function `syscall_context()` in `syscall.h` returns whether we are processing a syscall.
- There are two kinds of reading(or checking exactly). The function `check_user_string()` assumes that the user string is terminated with `'\0'`. It calls `is_user_addr()` and `get_user()` on user string byte after byte. If it doesn't detect an error before it encounters `'\0'`, user string is considered legal. Another function `check_user_string_bound()` is similar, except that `check_user_string_bound()` finishes checking when it meets length limitation without considering any byte value in user string.
- Only `sys_read()` requires writing to user address. So it makes user address check while writing to user buffer byte after byte.

B4: Suppose a system call causes a full page (4,096 bytes) of data to be copied from user space into the kernel. What is the least and the greatest possible number of inspections of the page table (e.g. calls to `pagedir_get_page()`) that might result? What about for a system call that only copies 2 bytes of data? Is there room for improvement in these numbers, and how much?

- In my implementation, the kernel will do this byte after byte. So in the first case, the answer is always 4096. In the second case, it's always 2.
- There is certainly room for improvement. If we know how many bytes to check, we can make this check one time per page. So in the first case, the answer is 1(least) and 2(greatest). The second case is the same.

B5: Briefly describe your implementation of the "wait" system call and how it interacts with process termination.

- Consider a situation that a parent thread P creates a child thread C through `exec()` and wait for C through `wait()`. When P calls `exec()`, it will call `thread_create()` to create thread C. In `thread_create()`, a child struct is initialized and inserted into P's `child_list`. In child struct, a semaphore named `terminated_sema` is initialized to value 0, which will play a key role in the wait process. In a call to `wait()` from P, P tries to down the `terminated_sema` related to C. However, the semaphore would only be upped when C calls `exit` (willingly or unwillingly), where is also the point to save C's exit status in its child struct. So P won't return from `sema_down()` until C exits. This is exactly what we expect.

B6: Any access to user program memory at a user-specified address can fail due to a bad pointer value. Such accesses must cause the process to be terminated. System calls are fraught with such accesses, e.g. a "write" system call requires reading the system call number from the user stack, then each of the call's three arguments, then an arbitrary amount of user memory, and any of these can fail at any point. This poses a design and error-handling problem: how do you best avoid obscuring the primary function of code in a morass of error-handling? Furthermore, when an error is detected, how do you ensure that all temporarily allocated resources (locks, buffers, etc.) are freed? In a few paragraphs, describe the strategy or strategies you adopted for managing these issues. Give an example.

- The general idea of error handling is that errors caused by bad pointers will terminate the process, and other types of errors will return an error code -1 to the calling process. Errors caused by bad pointers include failure to get syscall number, failure to get syscall arguments and bad addresses provided by user process. For example, a "write" syscall will be considered to cause a bad pointer error if it fails at the points mentioned above. However, the kernel will only return error code -1 if the failure is caused by a bad file descriptor.
- Checking syscall number and arguments are completed uniformly in `syscall_handler()` through `get_arg_n()` and `get_args()`. Then `syscall_handler()` calls specific handler according to syscall number. All handlers other than `sys_read()` and `sys_write()` are quite direct. These handlers just get the arguments from `argv[]`, call for checks on user addresses or file descriptors, then call the corresponding interfaces or handle the error according to the check result. Read and write are a bit more complex, but they follow the same error handling idea. It should be easy to understand following the comments in the function.
- Now let's talk about how to make sure all resources are released. Locks are only used on routines with a single entry and a single exit. So all locks acquired at the entry will be released at the exit. I only use a

buffer in `sys_read()` and it will be freed on every way out of the function.

## SYNCHRONIZATION

B7: The "exec" system call returns -1 if loading the new executable fails, so it cannot return before the new executable has completed loading. How does your code ensure this? How is the load success/failure status passed back to the thread that calls "exec"?

- The "exec" mainly calls `process_execute()`. All interesting things are in this function. Before a calling to `thread_create()`, parent thread P will initialize a semaphore called `child_load_sema` to value 0 in its thread struct. Then P calls `thread_create()` to create child thread C and may be blocked at a down operation on `child_load_sema`. Then C gets to run. When C finishes its loading, it will save loading result at `child_load_success` in P's thread struct and call `sema_up` on P's `child_load_sema` to notify it parent. Finally, P gets control again and it can get the load success/failure status from its `child_load_success`.

B8: Consider parent process P with child process C. How do you ensure proper synchronization and avoid race conditions when P calls `wait(C)` before C exits? After C exits? How do you ensure that all resources are freed in each case? How about when P terminates without waiting, before C exits? After C exits? Are there any special cases?

- First, let's talk about freeing resources. If P calls `wait()` on C, things we talked about in B5 ensure `wait()` will return at right time. And the child struct will be freed after P finishes waiting, which ensures the second call `wait()` on child C will return -1. If P doesn't call `wait()` on C, any remaining child struct will be freed when P exits through `free_child_list()`. So all resources are freed when P exits.
- Then I will talk about synchronizations. The child struct of C is shared between P and C. This struct is created and freed by P, but C can modify it. Only P has the right to access its child list. C can modify the child struct through a pointer called `child_msg` in its thread struct. But C has to make sure its parents is alive when it wants to make some modifications (otherwise may visit some freed memory). C gets this information through `is_parent_died` in C's thread struct. P will set C's `is_parent_died` when P calls `free_child_list()`. Another problem comes here. P must make sure C is alive when P tries to set C's `is_parent_died`. So P will check the child struct (owned by P) to know whether C is alive. Well, it's indeed a little complicated...
- In short, both P and C should know whether the other one is alive when they try to modify the other's data. P will check child status and decide whether to set C's `is_parent_alive` in `free_child_list()`. C will check `is_parent_died` and decide whether to modify the child struct. The key to avoiding race is to check the conditions and make corresponding modifications without interruption. So I use a global lock named `exit_lock` to make a part of `_exit()` function mutually exclusive. P's check-and-modify and C's check-and-modify both happen in this part.
- Now think about some cases mentioned above. If P waits for C, C exits before P without any race. However, now it's the case P doesn't wait for C. Our `exit_lock` ensures us to take `_exit()` as somehow an atom routine. If C exits before P, C is safe to modify its child struct because P hasn't free it. P won't try to modify C's `is_parent_died` because P knows C has exited. On the other side, P exits before C. P is safe to set C's `is_parent_died`. C won't try to modify its child struct because P has set C's `is_parent_died`.
- One special case is that loading failed when P call `exec()` to create C. But `thread_create()` has create C's child struct in P's `child_list`. In this case, C should notify P that C's thread is terminated in `start_process()`.

## RATIONALE

B9: Why did you choose to implement access to user memory from the kernel in the way that you did?

- I choose the second way mentioned in the Pintos documentation. This technique is normally faster because it takes advantage of the processor's MMU. And it can perform read and write operations while checking the address.

B10: What advantages or disadvantages can you see to your design for file descriptors?

- File descriptors are unique just within a single process. It's more convenient for a single process to manage its open file in an array. Since I use a int as a fd\_set, the maximum open files is 32. Although it seems enough in pintos, maximum 32 open files may not satisfy some big processes.

B11: The default tid\_t to pid\_t mapping is the identity mapping. If you changed it, what advantages are there to your approach?

- I didn't change it.